

Search Trees

- ◆ Explain Binary Search Trees
- ◆ Implement insert/delete operations
- ◆ Analyze the performance of binary search trees
- ◆ Explain balanced search trees

Maps and Hashtables - Review

- ◆ A map is a collection of key/value entries – keys are unique
- ◆ Map ADT: `get(k)`, `put(k, v)`, `remove(k)`, `size()`, `isEmpty()`, `entrySet()`, `keySet()`, `values()`
- ◆ We can use a doubly linked list to implement an unsorted map:
 - `put` takes $O(1)$ time
 - `get` and `remove` take $O(n)$ time
- ◆ Can we do better?
 - Use hash tables

Maps and Hashtables - Review

- ◆ A **hash function** h maps keys of a given type to integers in a fixed interval $[0, N - 1]$
- ◆ The integer $h(x)$ is called the **hash value** of key x
- ◆ A **hash table** for a given key type consists of
 - Hash function h
 - Array (called table) of size N
- ◆ When implementing a map with a hash table, the goal is to store item (k, o) at index $i = h(k)$
- ◆ A hash function composed of two functions:

Hash code:

$h_1: \text{keys} \rightarrow \text{integers}$

Compression function:

$h_2: \text{integers} \rightarrow [0, N - 1]$

Maps and Hashtables - Review

◆ Hash code implementation:

- Memory address of the key
- Integer cast - reinterpret the bits of the key as an integer
- Component sum - partition the bits of the key into components of fixed length (e.g., 16 or 32 bits) and we sum the components
- Polynomial accumulation

◆ Compression functions:


- Division: $h_2(y) = y \bmod N$
- Multiply, Add and Divide (MAD): $h_2(y) = (ay + b) \bmod N$


◆ Collision Handling

- Chaining
- Linear and quadratic probing

Presentation for use with the textbook **Data Structures and Algorithms in Java, 6th edition**, by M. T. Goodrich, R. Tamassia, and M. H. Goldwasser, Wiley, 2014


Binary Search Trees






Ordered Maps

- ◆ In a sorted map, entries are stored in order by their keys




- ◆ This allows us to support nearest neighbor queries:
 - ◆ Item with largest key less than or equal to k
 - ◆ Item with smallest key greater than or equal to k

Binary Search



- ◆ Binary search can perform nearest neighbor queries on an ordered map that is implemented with an array, sorted by key
 - similar to the high-low children's game
 - at each step, the number of candidate items is halved
 - terminates after $O(\log n)$ steps
- ◆ Example: **find**(7)




Binary Search Trees

- ◆ We can use a **search-tree structure** to efficiently implement a *sorted map*
- ◆ Binary trees are an excellent data structure for storing entries of a map, assuming we have an order relation defined on the keys.
- ◆ A **binary search tree** is a *proper binary tree*: each internal position p stores a key-value pair (k, v) such that:
 - Keys stored in the **left subtree** of p are **less** than k .
 - Keys stored in the **right subtree** of p are **greater** than k .

Binary Search Trees

- ◆ In the example below, the leaves of the tree serve only as “placeholders.”




- ◆ They can be represented as **null** references in practice, thereby reducing the number of nodes in half

Search Tables



- ◆ **A search table is an ordered map** implemented by means of a sorted sequence
 - We store the items in an array-based sequence, sorted by key
 - We use an external comparator for the keys
- ◆ **Performance:**
 - Searches take $O(\log n)$ time, using binary search
 - Inserting a new item takes $O(n)$ time, since in the worst case we have to shift $n/2$ items to make room for the new item
 - Removing an item takes $O(n)$ time, since in the worst case we have to shift $n/2$ items to compact the items after the removal
- ◆ The lookup table is **effective only for ordered maps of small size** or for maps on which searches are the most common operations, while insertions and removals are rarely performed (e.g., credit card authorizations)

Binary Search Trees




◆ A binary search tree is a binary tree storing keys (or key-value entries) at its internal nodes and satisfying the following property:

- Let u , v , and w be three nodes such that u is in the left subtree of v and w is in the right subtree of v . We have $key(u) \leq key(v) \leq key(w)$

◆ External nodes do not store items


◆ An **inorder** traversal of a binary search trees **visits the keys in increasing order**



Search

- ◆ To search for a key k , we trace a downward path starting at the root
- ◆ The next node visited depends on the comparison of k with the key of the current node
- ◆ If we reach a leaf, the key is not found
- ◆ Example: **get(4)**:
 - Call `TreeSearch(4, root)`
- ◆ The algorithms for nearest neighbor queries are similar

```
Algorithm TreeSearch( $k, v$ )  
  if T.isExternal ( $v$ )  
    return  $v$   
  if  $k < \text{key}(v)$   
    return TreeSearch( $k, \text{left}(v)$ )  
  else if  $k = \text{key}(v)$   
    return  $v$   
  else {  $k > \text{key}(v)$  }  
    return TreeSearch( $k, \text{right}(v)$ )
```



Insertion

- ◆ The map operation $\text{put}(k, v)$ begins with a search for an entry with key k .
 - If found, that entry's existing value is reassigned.
 - Otherwise, the **new entry can be inserted into the underlying tree** by expanding the leaf that was reached at the end of the failed search into an internal node.

Algorithm $\text{TreeInsert}(k, v)$:

Input: A search key k to be associated with value v

$p = \text{TreeSearch}(\text{root}(), k)$

if $k == \text{key}(p)$ **then**

 Change p 's value to (v)



else

$\text{expandExternal}(p, (k, v))$

- ◆ $\text{expandExternal}(p, e)$: Stores entry e at the external position p , and expands p to be internal, having two new leaves as children.

Insertion

- ◆ To perform operation $\text{put}(\mathbf{k}, o)$, we search for key \mathbf{k} (using `TreeSearch`)
- ◆ Assume \mathbf{k} is not already in the tree, and let \mathbf{w} be the leaf reached by the search
- ◆ We insert \mathbf{k} at node \mathbf{w} and expand \mathbf{w} into an internal node
- ◆ Example: insert 5





Deletion

- ◆ Deleting an entry from a binary search tree is a bit more complex than inserting a new entry
 - the position of an entry to be deleted might be anywhere in the tree (as opposed to insertions, which always occur at a leaf).
- ◆ To delete an entry with key k , we begin by calling `TreeSearch(root(), k)` to find the position p storing an entry with key equal to k (if any).

Deletion

- ◆ To perform operation **remove**(k), we search for key k
- ◆ Assume key k is in the tree, and let v be the node storing k
- ◆ If node v has a leaf child w , we remove v and w from the tree with operation **removeExternal**(w), which removes w and its parent
- ◆ Example: remove 4





Deletion (cont.)

- ◆ We consider the case where the key k to be removed is stored at a node v whose children are both internal



- we find the internal node w that follows v in an inorder traversal
- we copy $key(w)$ into node v
- we remove node w and its left child z (which must be a leaf) by means of operation `removeExternal(z)`

- ◆ Example: remove 3



Performance

- ◆ Consider an ordered map with n items implemented by means of a binary search tree of height h
 - the space used is $O(n)$
 - methods **get**, **put** and **remove** take $O(h)$ time
- ◆ The height h is $O(n)$ in the worst case and $O(\log n)$ in the best case



Performance

- ◆ on average, a binary search tree with n keys generated from a random series of insertions and removals of keys has expected height $O(\log n)$

Method	Running Time
size, isEmpty	$O(1)$
get, put, remove	$O(h)$
firstEntry, lastEntry	$O(h)$
ceilingEntry, floorEntry, lowerEntry, higherEntry	$O(h)$
subMap	$O(s + h)$
entrySet, keySet, values	$O(n)$

Java Implementation

- ◆ We define a **TreeMap** class that implements the sorted map ADT while using a binary search tree for storage.
- ◆ The **TreeMap** class is declared as a child of the **AbstractSortedMap** base class, thereby inheriting support for performing comparisons based upon a given (or default) **Comparator**, a nested **MapEntry** class for storing key-value pairs, and concrete implementations of methods **keySet** and values based upon the **entrySet** method, which we will provide.

Java Implementation


- ◆ For representing the tree structure, our TreeMap class maintains an instance of a subclass of the **LinkedBinaryTree** class from Section 8.3.1.
- ◆ In this implementation, we choose to represent the search tree as a *proper* binary tree, with explicit leaf nodes in the binary tree as sentinels, and map entries stored only at internal nodes.

Balanced Search Trees

- ◆ if we could assume a random series of insertions and removals, the standard binary search tree supports $O(\log n)$ expected running times for the basic map operations.
- ◆ However, we may only claim $O(n)$ worst-case time, because **some sequences of operations may lead to an unbalanced tree with height proportional to n .**
- ◆ A **balanced binary search tree** is a tree that automatically **keeps its height small** (guaranteed to be logarithmic) for a sequence of insertions and deletions.

Balanced Search Trees


- ◆ The primary operation to rebalance a binary search tree is known as a ***rotation***.
- ◆ During a rotation, we “rotate” a child to be above its parent.



- ◆ In the context of a tree-balancing algorithm, a rotation allows the shape of a tree to be modified while maintaining the search-tree property.

Balanced Search Trees

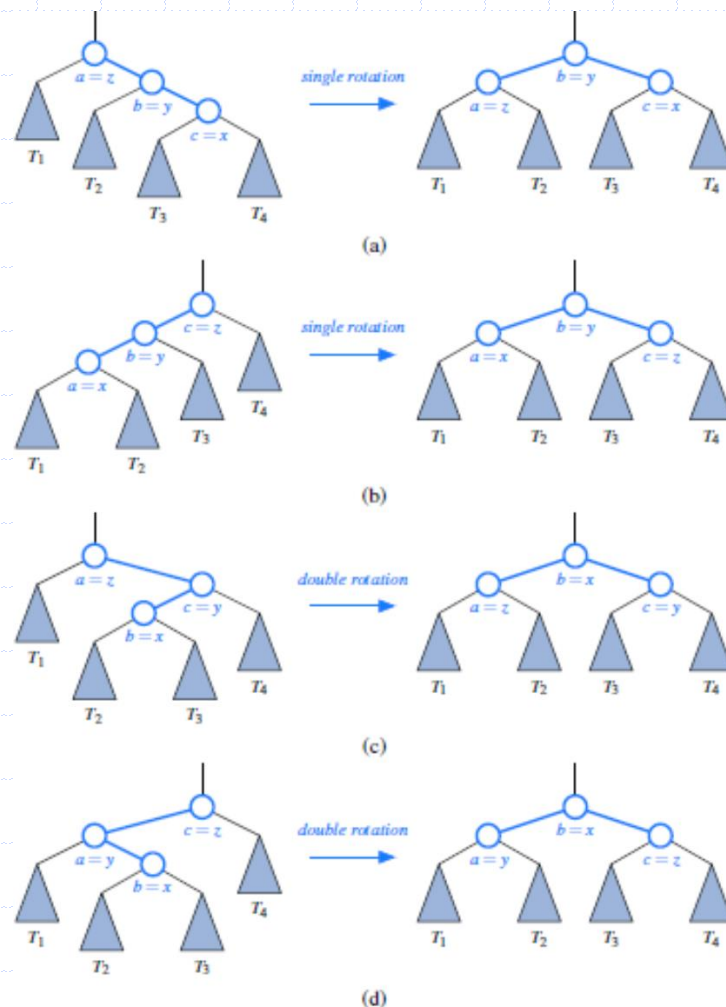
- ◆ A rotation can be performed to transform the left formation into the right, or the right formation into the left.
- ◆ Note that all keys in subtree T_1 have keys less than that of position x , all keys in subtree T_2 have keys that are between those of positions x and y , and all keys in subtree T_3 have keys that are greater than that of position y .



- ◆ In the first configuration, T_2 is the right subtree of position x ; in the second configuration, it is the left subtree of position y .

Balanced Search Trees

- ◆ One or more rotations can be combined to provide broader rebalancing within a tree.
- ◆ One such compound operation we consider is a *trinode restructuring*.



Java Framework for Balancing Search Trees

- ◆ The **TreeMap** class is designed in a way that allows it to be easily extended to provide more advanced tree-balancing strategies.
- ◆ Hooks for Rebalancing Operations
 - **rebalanceInsert(p)** is made from within the put method, after a new node is added to the tree at position p
 - **rebalanceDelete(p)** is made from within the remove method, after a node is deleted from the tree
 - **rebalanceAccess(p)** is made by any call to get, put, or remove that does not result in a structural change
- ◆ TreeMap class relies on storing the tree as an instance of a new nested class, **BalanceableBinaryTree**.
- ◆ That class is a specialization of the original **LinkedBinaryTree** class.