

Search Trees

- ◆ Explain Binary Search Trees
- ◆ Implement insert/delete operations
- ◆ Analyze the performance of binary search trees
- ◆ Explain balanced search trees

Maps and Hashtables - Review

- ◆ A map is a collection of key/value entries – keys are unique
- ◆ Map ADT: `get(k)`, `put(k, v)`, `remove(k)`, `size()`, `isEmpty()`, `entrySet()`, `keySet()`, `values()`
- ◆ We can use a **doubly linked list** to implement an unsorted map:
 - `put` takes $O(1)$ time
 - `get` and `remove` take $O(n)$ time
- ◆ Can we do better?
 - Use **hash tables**

Maps and Hashtables - Review

- ◆ A **hash function** h maps keys of a given type to integers in a fixed interval $[0, N - 1]$
- ◆ The integer $h(x)$ is called the **hash value** of key x
- ◆ A **hash table** for a given key type consists of
 - Hash **function** h
 - Array (called **table**) of size N
- ◆ When implementing a map with a hash table, the goal is to store item (k, o) at index $i = h(k)$
- ◆ A hash function composed of two functions:
 - Hash code:**
 $h_1: \text{keys} \rightarrow \text{integers}$
 - Compression function:**
 $h_2: \text{integers} \rightarrow [0, N - 1]$

Maps and Hashtables - Review

◆ Hash code implementation:

- Memory address of the key
- Integer cast - reinterpret the bits of the key as an integer
- Component sum - partition the bits of the key into components of fixed length (e.g., 16 or 32 bits) and we sum the components
- Polynomial accumulation

◆ Compression functions:

- Division: $h_2(y) = y \bmod N$
- Multiply, Add and Divide (MAD): $h_2(y) = (ay + b) \bmod N$


◆ Collision Handling

- Separate **Chaining**
- Linear and quadratic **probing**

◆ The **expected running time** of operations in a hash table is $O(1)$

Presentation for use with the textbook **Data Structures and Algorithms in Java, 6th edition**, by M. T. Goodrich, R. Tamassia, and M. H. Goldwasser, Wiley, 2014


Binary Search Trees



Binary Search Trees

- ◆ In this lecture we use a **search-tree structure** to efficiently implement a ***sorted map***.
- ◆ Recall that three most fundamental methods of a map are:
 - **get(k)**: Returns the value v associated with key k , if such an entry exists; otherwise returns null.
 - **put(k, v)**: Associates value v with key k , replacing and returning any existing value if the map already contains an entry with key equal to k .
 - **remove(k)**: Removes the entry with key equal to k , if one exists, and returns its value; otherwise returns null.

Ordered Maps




- ◆ In a **sorted map**, entries are stored in order by their keys:




- ◆ This allows us to **support nearest neighbor queries** (finding the item in a data structure that is closest to a given item, or key closest to lookup key):
 - ◆ Item with **largest key less** than or equal to k
 - ◆ Item with **smallest key greater** than or equal to k

Binary Search




- ◆ Binary search can perform nearest neighbor queries on an ordered map that is implemented with **an array, sorted by key**
 - similar to the high-low children's game (**l**-low, **m**-middle, **h**-high)
 - at each step, the number of candidate items is halved
 - terminates after **$O(\log n)$** steps
- ◆ Example: **find(7)**




Recall Binary Trees

- ◆ A **full binary tree** (sometimes **proper binary tree**) is a **tree** in which every node other than the leaves has two children.




- ◆ A complete set of
every letter
and all n




Binary Search Trees

- ◆ Binary trees are an **excellent data structure for storing entries of a map**, assuming we have an order relation defined on the keys.
- ◆ A **binary search tree** is a *proper binary tree*: each internal position p stores a key-value pair (k, v) such that:
 - Keys stored in the **left subtree** of p are **less** than k .
 - Keys stored in the **right subtree** of p are **greater** than k .



Traversing Binary Trees


- ◆ Traversal strategy specifies the order in which the current node, the left subtree, and the right subtree are visited.
 - **Inorder** traversal: the **left subtree** is visited first, then the **current node** followed by the **right subtree**
 - We assume the left subtree always comes before the right subtree



- The **Inorder** traversal gives: **1, 2, 4, 6, 8, 9**

Binary Search Trees

- ◆ In the example below, the leaves of the tree serve only as “placeholders.”



- ◆ They can be represented as **null** references in practice, thereby **reducing the number of nodes in half** (number of leaves in a full binary tree with n nodes is equal to $(n+1)/2$)


Search Tables vs Search Trees

- ◆ **A search table is an ordered map** implemented by means of a sorted sequence
 - We store the items in **an array-based sequence, sorted by key**
 - We use an **external comparator for the keys**
- ◆ **Performance:**
 - **Searches** take $O(\log n)$ time, **using binary search**
 - **Inserting** a new item takes $O(n)$ time, since in the worst case we have to shift n items to make room for the new item
 - **Removing** an item takes $O(n)$ time, since in the worst case we have to shift $n-1$ items to compact the items after the removal

Search Tables vs Search Trees

- ◆ The lookup table is **effective only for ordered maps of small size** or for maps on which searches are the most common operations, while insertions and removals are rarely performed (e.g., credit card authorizations)
- ◆ We can use a **search-tree structure** to efficiently implement a *sorted map*
 - *Achieve* **better performance for insertions and removals**

Binary Search Trees




◆ A binary search tree is a binary tree **storing** keys (or key-value entries) **at its internal nodes** and satisfying the following property:

- Let u , v , and w be three nodes such that u is in the left subtree of v and w is in the right subtree of v . We have $key(u) \leq key(v) \leq key(w)$

◆ External nodes do not store items

◆ An **inorder** traversal of a binary search trees **visits the keys in increasing order**


- The left-most child has the smallest key
- The right-most child has the largest key



Search

- ◆ To search for a key k , we **trace a downward path starting at the root**
- ◆ The next node visited depends on the **comparison of k with the key of the current node**
- ◆ If we reach a leaf, the key is not found
- ◆ Example: **get(4)**:
 - Call **TreeSearch(4, root)**
- ◆ The algorithms for nearest neighbor queries are similar

```
Algorithm TreeSearch( $k, v$ )  
  if T.isExternal ( $v$ )  
    return  $v$   
  if  $k < \text{key}(v)$   
    return TreeSearch( $k, \text{left}(v)$ )  
  else if  $k = \text{key}(v)$   
    return  $v$   
  else {  $k > \text{key}(v)$  }  
    return TreeSearch( $k, \text{right}(v)$ )
```



Insertion

- ◆ The map operation **put**(k , v) begins with a search for an entry with key k .
 - If found, that entry's existing value is reassigned.
 - Otherwise, the **new entry can be inserted into the underlying tree** by **expanding the leaf that was reached at the end of the failed search into an internal node**.

Algorithm TreeInsert(k , v):

Input: A search key k to be associated with value v

$p = \text{TreeSearch}(\text{root}(), k)$

if $k == \text{key}(p)$ then

 Change p 's value to (v)



else

 expandExternal(p , (k , v))

- ◆ **expandExternal**(p , e): Stores entry e at the external position p , and **expands** p to be internal, having two new leaves as children.

Insertion

- ◆ To perform operation **put(\mathbf{k} , \mathbf{o})**, we search for key \mathbf{k} (using TreeSearch)
- ◆ Assume \mathbf{k} is not already in the tree, and let \mathbf{w} be the leaf reached by the search
- ◆ We insert \mathbf{k} at node \mathbf{w} and expand \mathbf{w} into an internal node
- ◆ Example: **insert 5**





Deletion

- ◆ Deleting an entry from a binary search tree is a bit more complex than inserting a new entry
 - the position of an entry to be deleted might be anywhere in the tree (as opposed to **insertions, which always occur at a leaf**).
 - ◆ node has **no children**
 - ◆ node has **one child**
 - ◆ node is **internal**
- ◆ To delete an entry with key k , we begin by calling `TreeSearch(root(), k)` to find the position p storing an entry with key equal to k (if any).


Deletion – node has no children

- ◆ Assume key k is in the tree, and let v be the node storing k
 - Remove v and its leafs (placeholders)
 - Replace v with a leaf (placeholder)
- ◆ Example: **remove 4**



Deletion – node has one child


- ◆ To perform operation **remove(k)**, we search for key k
- ◆ Assume key k is in the tree, and let v be the node storing k
- ◆ If **node v has a leaf child w** , we remove v and w from the tree with operation **removeExternal(w)**, which removes w and its parent
- ◆ Example: **remove 4**



Deletion – node has two children

- ◆ We consider the case where the key k to be removed is stored at a node v whose **children are both internal**
 - we find the internal node w that follows v in an **inorder traversal**
 - we **copy** $key(w)$ into node v
 - we **remove node w and its left child z** (which must be a leaf, why?) by means of operation **removeExternal(z)**

◆ Example: **remove 3**



-



Performance

- ◆ On average, a binary search tree with n keys generated from a random series of insertions and removals of keys has **expected height** $O(\log n)$

Method	Running Time
size, isEmpty	$O(1)$
get, put, remove	$O(h)$
firstEntry, lastEntry	$O(h)$
ceilingEntry, floorEntry, lowerEntry, higherEntry	$O(h)$
subMap	$O(s + h)$
entrySet, keySet, values	$O(n)$

Java Implementation

- ◆ We define a **TreeMap** class that **implements the sorted map** ADT while using a binary search tree for storage.
- ◆ The **TreeMap** class is declared as a child of the **AbstractSortedMap** base class, thereby inheriting support for performing comparisons based upon a given (or default) **Comparator**, a nested **MapEntry** class for storing key-value pairs, and concrete implementations of methods **keySet** and values based upon the **entrySet** method, which we will provide.

Java Implementation

- ◆ For representing the tree structure, our TreeMap class maintains an instance of a subclass of the **LinkedBinaryTree** class from Section 8.3.1.
- ◆ In this implementation, we choose to represent the **search tree as a *proper* binary tree**, with explicit leaf nodes in the binary tree as sentinels (**null**), and map entries stored only at internal nodes.

Java Implementation

```
1  /** An implementation of a sorted map using a binary search tree. */
2  public class TreeMap<K,V> extends AbstractSortedMap<K,V> {
3      // To represent the underlying tree structure, we use a specialized subclass of the
4      // LinkedBinaryTree class that we name BalanceableBinaryTree (see Section 11.2).
5      protected BalanceableBinaryTree<K,V> tree = new BalanceableBinaryTree<>();
6
7      /** Constructs an empty map using the natural ordering of keys. */
8      public TreeMap() {
9          super(); // the AbstractSortedMap constructor
10         tree.addRoot(null); // create a sentinel leaf as root
11     }
12     /** Constructs an empty map using the given comparator to order keys. */
13     public TreeMap(Comparator<K> comp) {
14         super(comp); // the AbstractSortedMap constructor
15         tree.addRoot(null); // create a sentinel leaf as root
16     }
17     /** Returns the number of entries in the map. */
18     public int size() {
19         return (tree.size() - 1) / 2; // only internal nodes have entries
20     }
21     /** Utility used when inserting a new entry at a leaf of the tree */
22     private void expandExternal(Position<Entry<K,V>> p, Entry<K,V> entry) {
23         tree.set(p, entry); // store new entry at p
24         tree.addLeft(p, null); // add new sentinel leaves as children
25         tree.addRight(p, null);
26     }
27 }
```

Java Implementation

```
28 // Omitted from this code fragment, but included in the online version of the code,
29 // are a series of protected methods that provide notational shorthands to wrap
30 // operations on the underlying linked binary tree. For example, we support the
31 // protected syntax root() as shorthand for tree.root() with the following utility:
32 protected Position<Entry<K,V>> root() { return tree.root(); }
33
34 /** Returns the position in p's subtree having given key (or else the terminal leaf).*/
35 private Position<Entry<K,V>> treeSearch(Position<Entry<K,V>> p, K key) {
36     if (isExternal(p))
37         return p; // key not found; return the final leaf
38     int comp = compare(key, p.getElement());
39     if (comp == 0)
40         return p; // key found; return its position
41     else if (comp < 0)
42         return treeSearch(left(p), key); // search left subtree
43     else
44         return treeSearch(right(p), key); // search right subtree
45 }
```

Code Fragment 11.3: Beginning of a TreeMap class based on a binary search tree.

Java Implementation

```
46  /** Returns the value associated with the specified key (or else null). */
47  public V get(K key) throws IllegalArgumentException {
48      checkKey(key); // may throw IllegalArgumentException
49      Position<Entry<K,V>> p = treeSearch(root(), key);
50      rebalanceAccess(p); // hook for balanced tree subclasses
51      if (isExternal(p)) return null; // unsuccessful search
52      return p.getElement().getValue(); // match found
53  }
54  /** Associates the given value with the given key, returning any overridden value.*/
55  public V put(K key, V value) throws IllegalArgumentException {
56      checkKey(key); // may throw IllegalArgumentException
57      Entry<K,V> newEntry = new MapEntry<>(key, value);
58      Position<Entry<K,V>> p = treeSearch(root(), key);
59      if (isExternal(p)) { // key is new
60          expandExternal(p, newEntry);
61          rebalanceInsert(p); // hook for balanced tree subclasses
62          return null;
63      } else { // replacing existing key
64          V old = p.getElement().getValue();
65          set(p, newEntry);
66          rebalanceAccess(p); // hook for balanced tree subclasses
67          return old;
68      }
69  }
```

Java Implementation

```
70  /** Removes the entry having key k (if any) and returns its associated value. */
71  public V remove(K key) throws IllegalArgumentException {
72      checkKey(key); // may throw IllegalArgumentException
73      Position<Entry<K,V>> p = treeSearch(root(), key);
74      if (isExternal(p)) { // key not found
75          rebalanceAccess(p); // hook for balanced tree subclasses
76          return null;
77      } else {
78          V old = p.getElement().getValue();
79          if (isInternal(left(p)) && isInternal(right(p))) { // both children are internal
80              Position<Entry<K,V>> replacement = treeMax(left(p));
81              set(p, replacement.getElement());
82              p = replacement;
83          } // now p has at most one child that is an internal node
84          Position<Entry<K,V>> leaf = (isExternal(left(p)) ? left(p) : right(p));
85          Position<Entry<K,V>> sib = sibling(leaf);
86          remove(leaf);
87          remove(p); // sib is promoted in p's place
88          rebalanceDelete(sib); // hook for balanced tree subclasses
89          return old;
90      }
91  }
```

Code Fragment 11.4: Primary map operations for the TreeMap class.

Java Implementation

```
92  /** Returns the position with the maximum key in subtree rooted at Position p. */
93  protected Position<Entry<K,V>> treeMax(Position<Entry<K,V>> p) {
94      Position<Entry<K,V>> walk = p;
95      while (isInternal(walk))
96          walk = right(walk);
97      return parent(walk);                // we want the parent of the leaf
98  }
99  /** Returns the entry having the greatest key (or null if map is empty). */
100 public Entry<K,V> lastEntry() {
101     if (isEmpty()) return null;
102     return treeMax(root()).getElement();
103 }
104 /** Returns the entry with greatest key less than or equal to given key (if any). */
105 public Entry<K,V> floorEntry(K key) throws IllegalArgumentException {
106     checkKey(key);                      // may throw IllegalArgumentException
107     Position<Entry<K,V>> p = treeSearch(root(), key);
108     if (isInternal(p)) return p.getElement(); // exact match
109     while (isRoot(p)) {
110         if (p == right(parent(p)))
111             return parent(p).getElement(); // parent has next lesser key
112         else
113             p = parent(p);
114     }
115     return null;                        // no such floor exists
116 }
```


Java Implementation

```
117  /** Returns the entry with greatest key strictly less than given key (if any). */
118  public Entry<K,V> lowerEntry(K key) throws IllegalArgumentException {
119      checkKey(key); // may throw IllegalArgumentException
120      Position<Entry<K,V>> p = treeSearch(root(), key);
121      if (isInternal(p) && isInternal(left(p)))
122          return treeMax(left(p)).getElement(); // this is the predecessor to p
123      // otherwise, we had failed search, or match with no left child
124      while (isRoot(p)) {
125          if (p == right(parent(p)))
126              return parent(p).getElement(); // parent has next lesser key
127          else
128              p = parent(p);
129      }
130      return null; // no such lesser key exists
131  }
```

Code Fragment 11.5: A sample of the sorted map operations for the `TreeMap` class. The symmetrical utility, `treeMin`, and public methods `firstEntry`, `ceilingEntry`, and `higherEntry` are available online.

Java Implementation

```
132  /** Returns an iterable collection of all key-value entries of the map. */
133  public Iterable<Entry<K,V>> entrySet() {
134      ArrayList<Entry<K,V>> buffer = new ArrayList<>(size());
135      for (Position<Entry<K,V>> p : tree.inorder())
136          if (isInternal(p)) buffer.add(p.getElement());
137      return buffer;
138  }
139  /** Returns an iterable of entries with keys in range [fromKey, toKey). */
140  public Iterable<Entry<K,V>> subMap(K fromKey, K toKey) {
141      ArrayList<Entry<K,V>> buffer = new ArrayList<>(size());
142      if (compare(fromKey, toKey) < 0) // ensure that fromKey < toKey
143          subMapRecurse(fromKey, toKey, root(), buffer);
144      return buffer;
145  }
146  private void subMapRecurse(K fromKey, K toKey, Position<Entry<K,V>> p,
147                          ArrayList<Entry<K,V>> buffer) {
148      if (isInternal(p))
149          if (compare(p.getElement(), fromKey) < 0)
150              // p's key is less than fromKey, so any relevant entries are to the right
151              subMapRecurse(fromKey, toKey, right(p), buffer);
152          else {
153              subMapRecurse(fromKey, toKey, left(p), buffer); // first consider left subtree
154              if (compare(p.getElement(), toKey) < 0) { // p is within range
155                  buffer.add(p.getElement()); // so add it to buffer, and consider
156                  subMapRecurse(fromKey, toKey, right(p), buffer); // right subtree as well
157              }
158          }
159  }
```


Code Fragment 11.6: TreeMap operations supporting iteration of the entire map, or a portion of the map with a given key range.

Balanced Search Trees

- ◆ if we could assume a **random series of insertions** and removals, the standard binary search tree supports **$O(\log n)$ expected running times** for the basic map operations.
- ◆ However, we may only claim $O(n)$ worst-case time, because **some sequences of operations may lead to an unbalanced tree with height proportional to n .**
- ◆ A **balanced binary search tree** is a tree that automatically **keeps its height small** (guaranteed to be logarithmic) for a sequence of insertions and deletions.

Balanced Search Trees


- ◆ The primary operation to rebalance a binary search tree is known as a ***rotation***.
- ◆ During a rotation, we “rotate” a child to be above its parent.



- ◆ In the context of a tree balancing algorithm, a rotation allows the shape of a tree to be modified while maintaining the search-tree property.

Balanced Search Trees

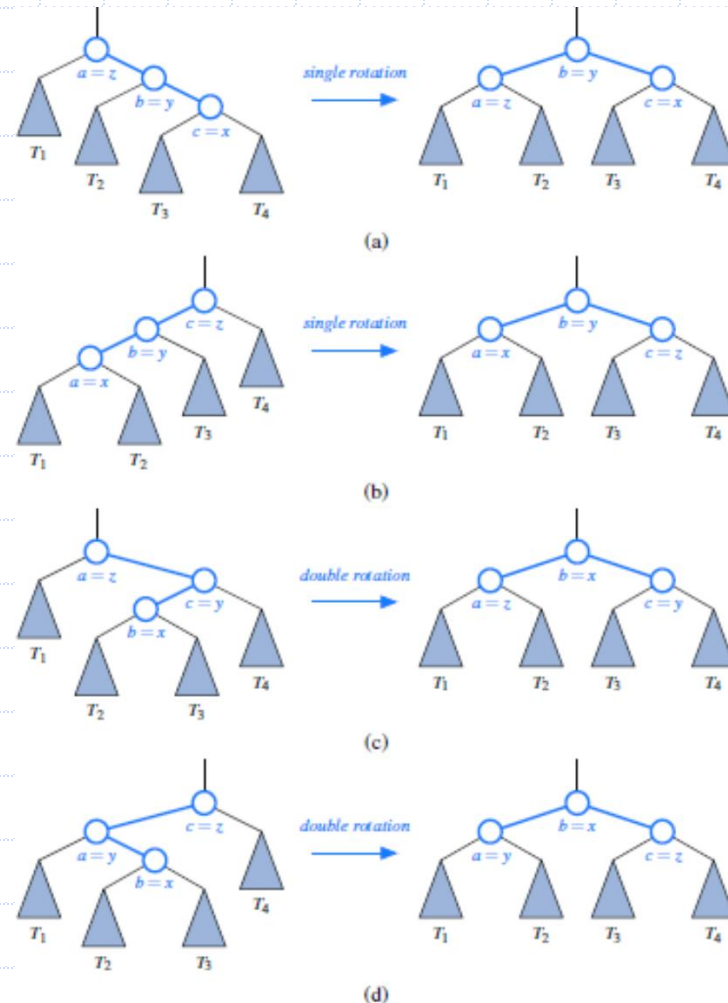
- ◆ A rotation can be performed to transform the left formation into the right, or the right formation into the left.
- ◆ Note that all keys in subtree T_1 have keys less than that of position x , all keys in subtree T_2 have keys that are between those of positions x and y , and all keys in subtree T_3 have keys that are greater than that of position y .



- ◆ In the first configuration, T_2 is the right subtree of position x ; in the second configuration, it is the left subtree of position y .

Balanced Search Trees

- ◆ One or more rotations can be combined to provide broader rebalancing within a tree.
- ◆ One such compound operation we consider is a **trinode restructuring**.



Java Framework for Balancing Search Trees

- ◆ The **TreeMap** class is designed in a way that allows it to be easily extended to provide more advanced tree-balancing strategies.
- ◆ Hooks for Rebalancing Operations
 - **rebalanceInsert(p)** is made from within the put method, after a new node is added to the tree at position p
 - **rebalanceDelete(p)** is made from within the remove method, after a node is deleted from the tree
 - **rebalanceAccess(p)** is made by any call to get, put, or remove that does not result in a structural change
- ◆ TreeMap class relies on storing the tree as an instance of a new nested class, **BalanceableBinaryTree**.
- ◆ That class is a specialization of the original **LinkedBinaryTree** class.