

Sorting

- Explain Merge-sort
- Explain quick-sort
- Analyze sorting algorithms
- Compare sorting algorithms


Search Trees - Review

- ◆ Use a **search-tree structure** to efficiently implement a *sorted map*
- ◆ A **binary search tree** is a *proper binary tree*: each internal position p stores a key-value pair (k, v) such that:
 - Keys stored in the **left subtree** of p are **less** than k .
 - Keys stored in the **right subtree** of p are **greater** than k .
- ◆ Operations:
 - Search
 - Insert
 - Delete

Search Trees - Review



- ◆ Search for a key:
- ◆ Example: `get(4)`:
 - Call `TreeSearch(4, root)`

```
Algorithm TreeSearch( $k, v$ )  
  if  $T.isExternal(v)$   
    return  $v$   
  if  $k < key(v)$   
    return TreeSearch( $k, left(v)$ )  
  else if  $k = key(v)$   
    return  $v$   
  else {  $k > key(v)$  }  
    return TreeSearch( $k, right(v)$ )
```





Search Trees - Review

- ◆ Insertion:
- ◆ To perform operation $\text{put}(\mathbf{k}, o)$, we search for key \mathbf{k} (using TreeSearch)
- ◆ Assume \mathbf{k} is not already in the tree, and let \mathbf{w} be the leaf reached by the search
- ◆ We insert \mathbf{k} at node \mathbf{w} and expand \mathbf{w} into an internal node
- ◆ Example: insert 5




Deletion – node has no children

- ◆ Assume key k is in the tree, and let v be the node storing k
 - *Remove v and its leafs (placeholders)*
 - *Replace v with a leaf (placeholder)*
- ◆ Example: **remove 4**



Deletion – node has one child

- ◆ To perform operation **remove(k)**, we search for key k
- ◆ Assume key k is in the tree, and let v be the node storing k
- ◆ If **node v has a leaf child w** , we remove v and w from the tree with operation **removeExternal(w)**, which removes w and its parent
- ◆ Example: **remove 4**





Deletion – node has two children

- ◆ We consider the case where the key k to be removed is stored at a node v whose **children are both internal**

- we find the internal node w that follows v in an **inorder traversal**
- we copy $key(w)$ into node v
- we remove node w and its left child z (which must be a leaf) by means of operation **removeExternal(z)**

- ◆ Example: **remove 3**




Performance

- ◆ On average, a binary search tree with n keys generated from a random series of insertions and removals of keys has expected height $O(\log n)$

Method	Running Time
size, isEmpty	$O(1)$
get, put, remove	$O(h)$
firstEntry, lastEntry	$O(h)$
ceilingEntry, floorEntry, lowerEntry, higherEntry	$O(h)$
subMap	$O(s + h)$
entrySet, keySet, values	$O(n)$

Presentation for use with the textbook **Data Structures and Algorithms in Java, 6th edition**, by M. T. Goodrich, R. Tamassia, and M. H. Goldwasser, Wiley, 2014

Merge Sort



Divide-and-Conquer

◆ **Divide-and conquer** is a general **algorithm design paradigm**:

- **Divide**: divide the input data S in **two disjoint subsets** S_1 and S_2
- **Recur**: **solve the subproblems** associated with S_1 and S_2
- **Conquer**: **combine the solutions** for S_1 and S_2 into a solution for S

◆ The **base case** for the recursion are subproblems of size 0 or 1

◆ **Merge-sort** is a sorting algorithm based on the divide-and-conquer paradigm

◆ Like heap-sort

- It has $O(n \log n)$ running time

◆ Unlike heap-sort

- It does not use an auxiliary priority queue
- It accesses data in a **sequential** manner (suitable to sort data on a disk)

Merge-Sort

- ◆ Merge-sort on an input sequence S with n elements consists of three steps:
 - **Divide**: partition S into two sequences S_1 and S_2 of about $n/2$ elements each
 - **Recur**: recursively sort S_1 and S_2
 - **Conquer**: merge sorted sequences S_1 and S_2 into a unique sorted sequence

Algorithm *mergeSort*(S)

Input sequence S with n elements

Output sequence S sorted according to C

if $S.size() > 1$

$(S_1, S_2) \leftarrow partition(S, n/2)$

mergeSort(S_1)

mergeSort(S_2)

$S \leftarrow merge(S_1, S_2)$

Merging Two Sorted Sequences

- ◆ The conquer step of merge-sort consists of **merging two sorted sequences** A and B into a sorted sequence S containing the **union** of the elements of A and B
- ◆ Merging two sorted sequences, each with $n/2$ elements and **implemented by means of a doubly linked list**, takes $O(n)$ time

Algorithm *merge*(A, B)

Input sequences A and B with $n/2$ elements each

Output sorted sequence of $A \cup B$ in increasing order

$S \leftarrow$ empty sequence

while $\neg A.isEmpty() \wedge \neg B.isEmpty()$

if $A.first().element() < B.first().element()$

$S.addLast(A.remove(A.first()))$

else

$S.addLast(B.remove(B.first()))$

while $\neg A.isEmpty()$

$S.addLast(A.remove(A.first()))$



while $\neg B.isEmpty()$

$S.addLast(B.remove(B.first()))$

return S

Java Merge Implementation

```
1  /** Merge contents of arrays S1 and S2 into properly sized array S. */
2  public static <K> void merge(K[ ] S1, K[ ] S2, K[ ] S, Comparator<K> comp) {
3      int i = 0, j = 0;
4      while (i + j < S.length) {
5          if (j == S2.length || (i < S1.length && comp.compare(S1[i], S2[j]) < 0))
6              S[i+j] = S1[i++];           // copy ith element of S1 and increment i
7          else
8              S[i+j] = S2[j++];           // copy jth element of S2 and increment j
9      }
10 }
```




Java Merge Implementation

◆ Picture explanation:

- $i=3, j=2$
- $S1[3] > S2[2] // 11 > 10$
- $S[i+j] = S[5] = S2[2] // 10$

◆ Code explanation:

- if $j == S2.length$ **or** ($i < S1.length$ and $S1[i] < S2[j]$) copy element from S1
- Otherwise:
 - ◆ Otherwise, copy element from S2




Java Merge-Sort Implementation

```
1  /** Merge-sort contents of array S. */
2  public static <K> void mergeSort(K[ ] S, Comparator<K> comp) {
3      int n = S.length;
4      if (n < 2) return;           // array is trivially sorted
5      // divide
6      int mid = n/2;
7      K[ ] S1 = Arrays.copyOfRange(S, 0, mid);    // copy of first half
8      K[ ] S2 = Arrays.copyOfRange(S, mid, n);    // copy of second half
9      // conquer (with recursion)
10     mergeSort(S1, comp);           // sort copy of first half
11     mergeSort(S2, comp);           // sort copy of second half
12     // merge results
13     merge(S1, S2, S, comp);        // merge sorted halves back into original
14 }
```


Merge-Sort Tree

- ◆ An execution of merge-sort is depicted by a binary tree
 - each node represents a recursive call of merge-sort and stores
 - ◆ unsorted sequence before the execution and its partition
 - ◆ sorted sequence at the end of the execution
 - the root is the initial call
 - the leaves are calls on subsequences of size 0 or 1




Execution Example

◆ Partition




Execution Example (cont.)

◆ Recursive call, partition




Execution Example (cont.)

◆ Recursive call, partition




Execution Example (cont.)

- ◆ Recursive call, reaches base case




Execution Example (cont.)

◆ Recursive call, reaches base case



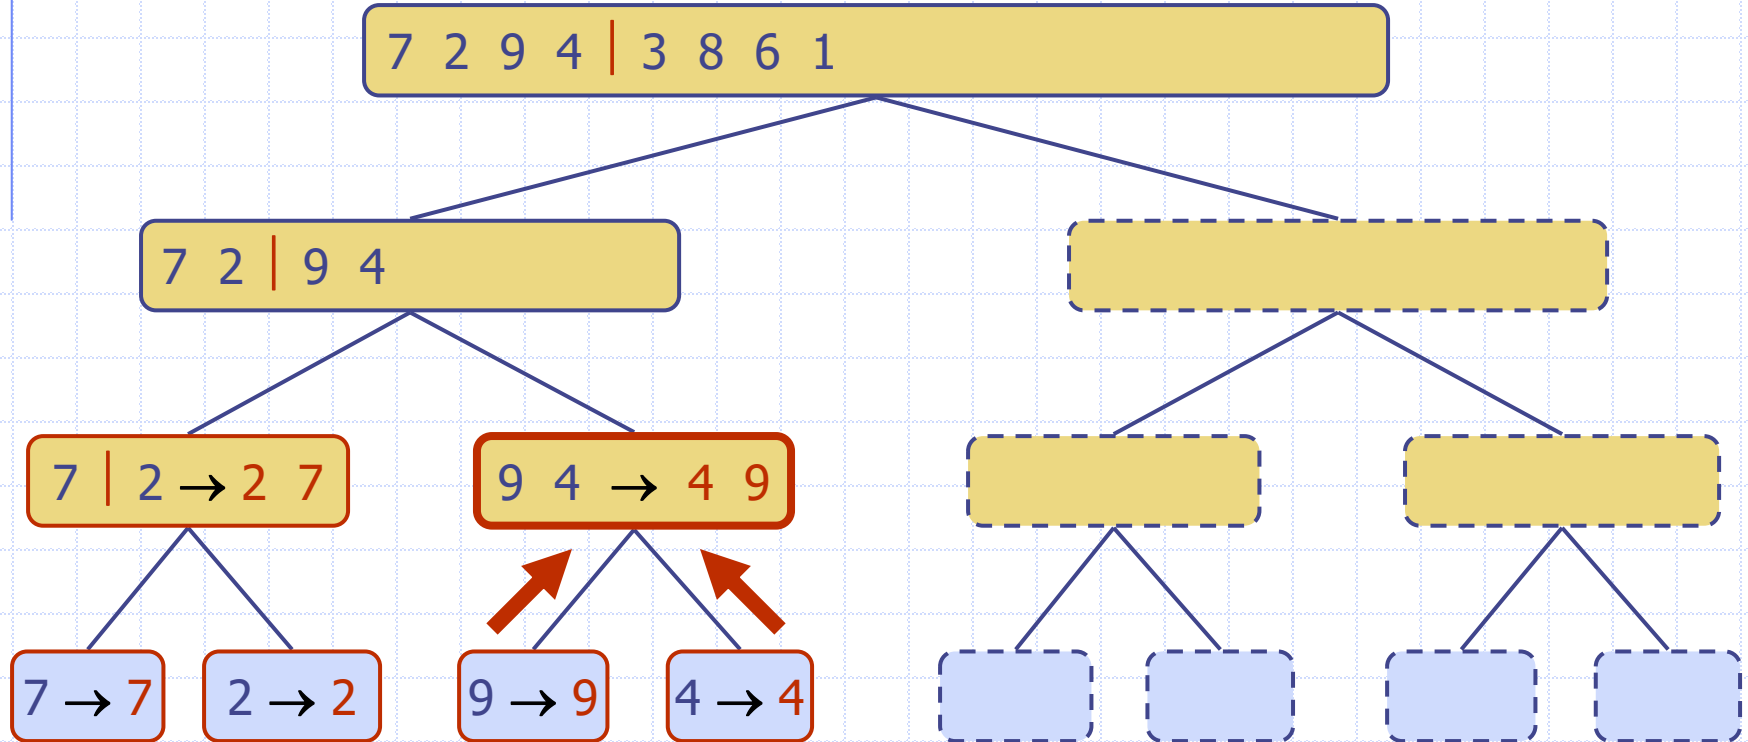
Execution Example (cont.)

◆ Merge



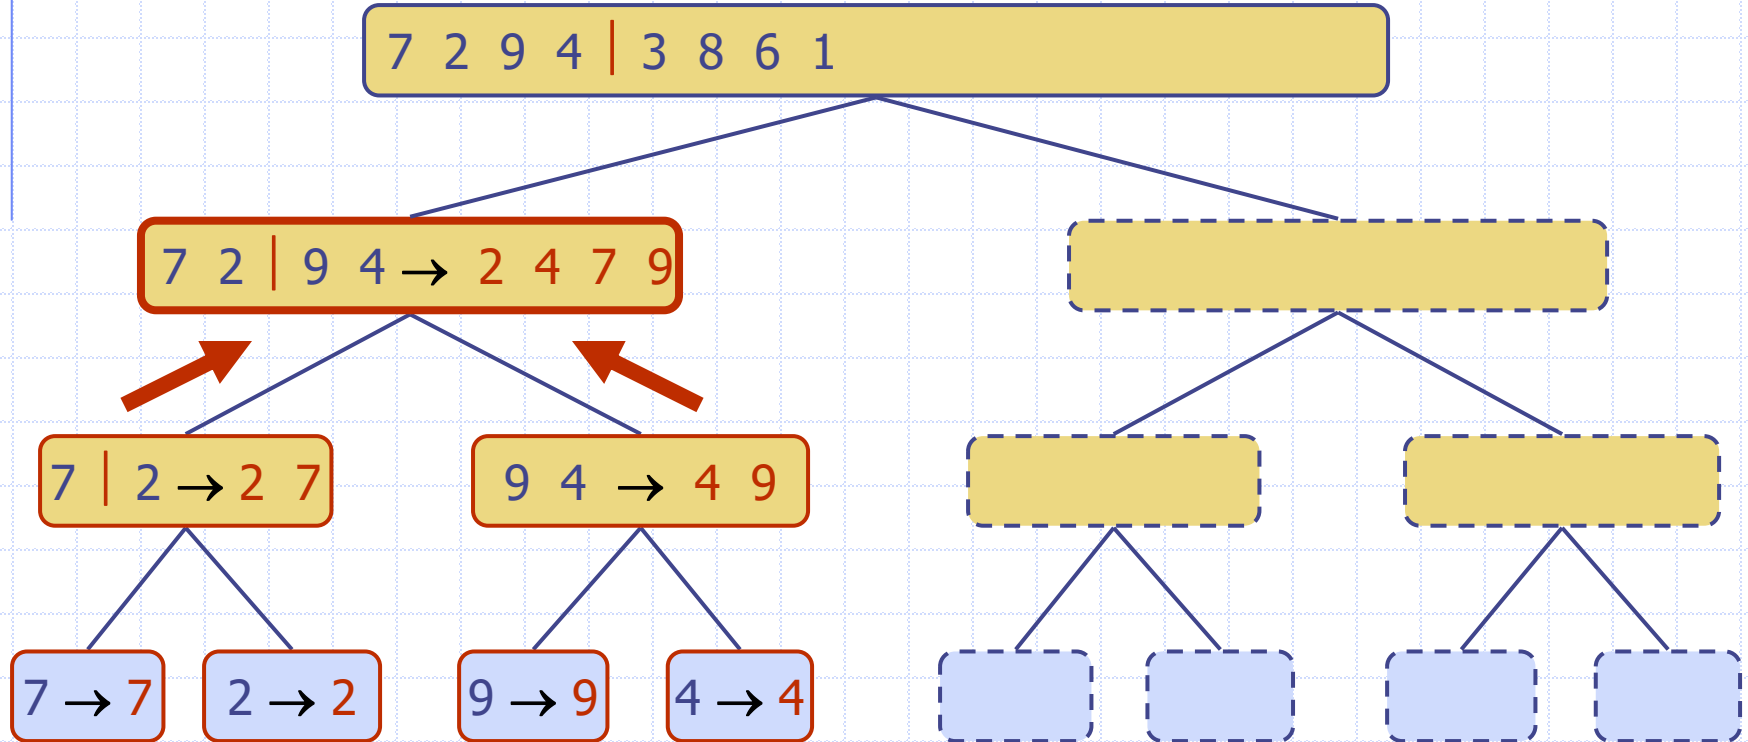
Execution Example (cont.)

- ◆ Recursive call, ..., base case, merge



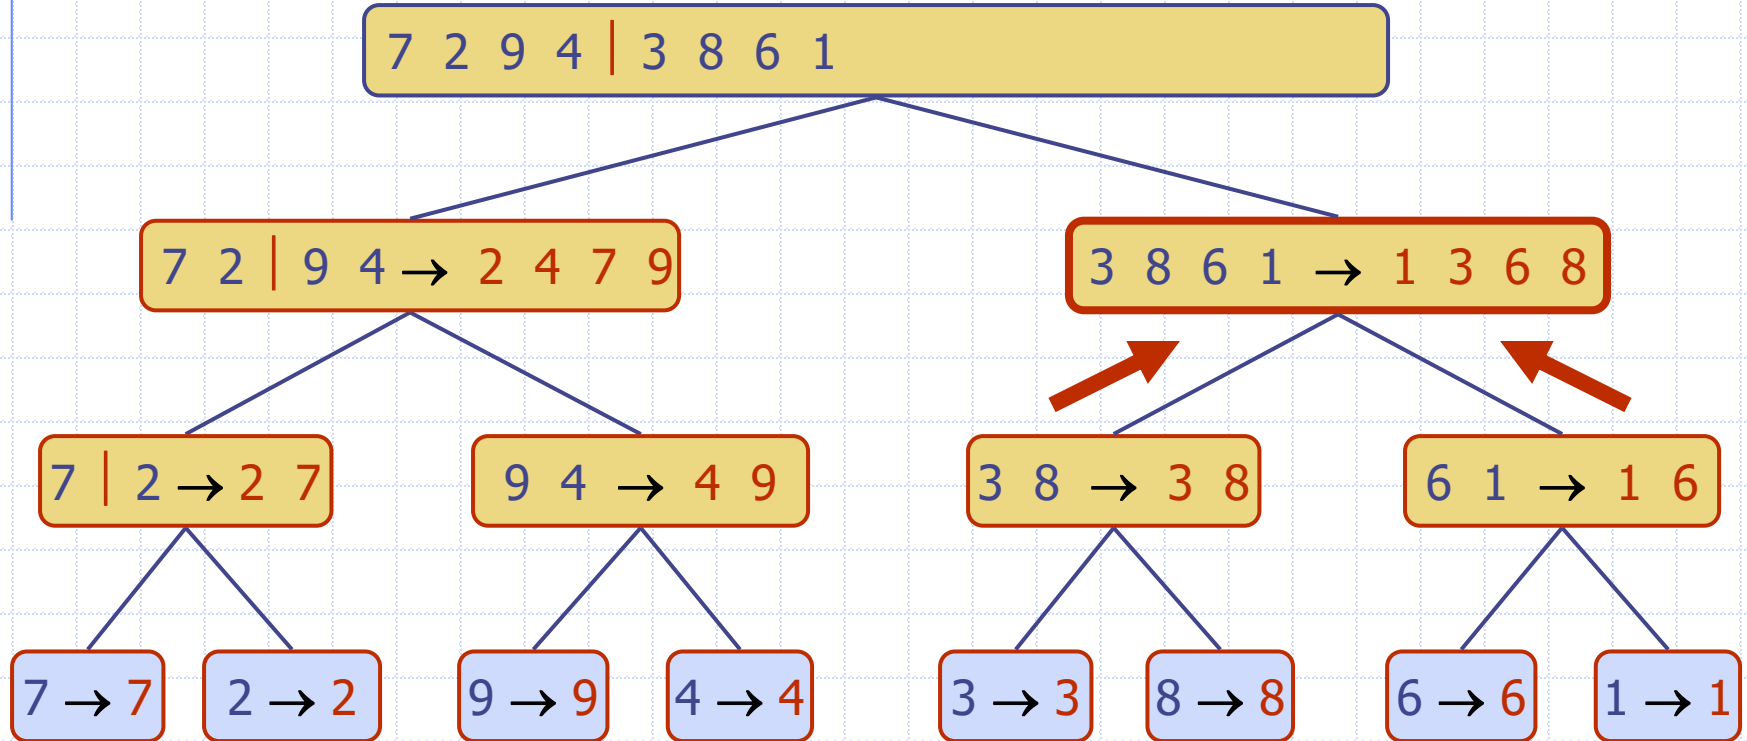
Execution Example (cont.)

◆ Merge




Execution Example (cont.)

◆ Recursive call, ..., merge, merge



Execution Example (cont.)

◆ Merge



Analysis of Merge-Sort

- ◆ Running time of the **merge** algorithm:
 - Let **n_1** and **n_2** be the number of elements of S_1 and S_2 , respectively.
 - It is clear that the operations (addLast, removeFirst) performed inside each pass of the while loop take $O(1)$ time.
 - The key observation is that during each iteration of the loop, one element is copied from either S_1 or S_2 into S (and that element is considered no further).
 - Therefore, the number of iterations of the loop is n_1+n_2 .
- ◆ Thus, the running time of algorithm **merge** is **$O(n_1+n_2)$** , **hence linear in the size of merged sequence.**

Analysis of Merge-Sort

- ◆ Running time of the entire merge-sort algorithm:
 - We account for the **amount of time spent within each recursive call**, but excluding any time spent waiting for successive recursive calls to terminate.
 - In the case of our mergeSort method, we account for the **time to divide** the sequence into two subsequences, and the **call to merge** to combine the two sorted sequences, but we exclude the two recursive calls to mergeSort.
 - We use a **merge-sort tree** T to guide our analysis.

Analysis of Merge-Sort

- ◆ Consider a recursive call associated with a node v of the merge-sort tree T :
 - the **divide step** at node v runs in time proportional to the size of the sequence for v .
 - the **merging step** also takes time that is linear in the size of the merged sequence
 - the **size of the sequence** handled by the recursive call associated with v is equal to $n/2^i$


depth #seqs size

0 1 n

1 2 $n/2^1$

i 2^i $n/2^i$

...



Analysis of Merge-Sort

- ◆ The running time of merge-sort is equal to the sum of the times spent at the nodes of T .
 - *recall that T has exactly 2^i nodes at depth i .*
 - the overall time spent at all the nodes of T at depth i is $O(2^i \cdot n/2^i)$, which is **$O(n)$** .
 - Recall that the height h of the merge-sort tree is **$O(\log n)$** .
 - Therefore, algorithm merge-sort sorts a sequence S of size n in **$O(n \log n)$** time, assuming two elements of S can be compared in $O(1)$ time.

Analysis of Merge-Sort

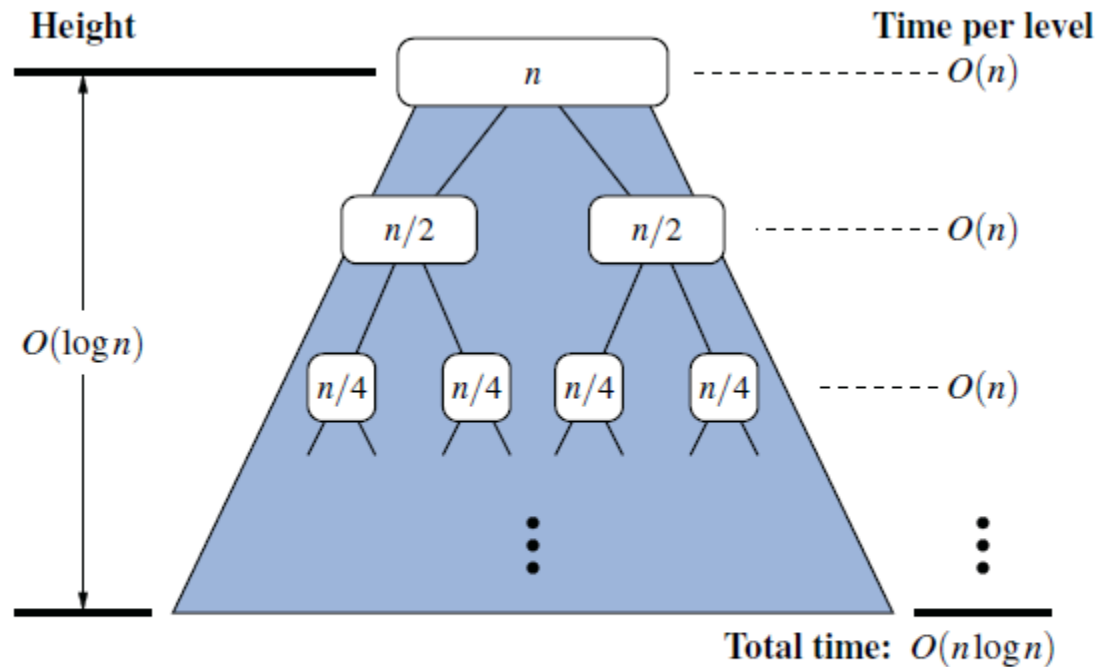



Figure 12.6: A visual analysis of the running time of merge-sort. Each node represents the time spent in a particular recursive call, labeled with the size of its subproblem.

Summary of Sorting Algorithms

Algorithm	Time	Notes
selection-sort	$O(n^2)$	<ul style="list-style-type: none">▪ slow▪ in-place▪ for small data sets (< 1K)
insertion-sort	$O(n^2)$	<ul style="list-style-type: none">▪ slow▪ in-place▪ for small data sets (< 1K)
heap-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ fast▪ in-place▪ for large data sets (1K — 1M)
merge-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ fast▪ sequential data access▪ for huge data sets (> 1M)


Presentation for use with the textbook **Data Structures and Algorithms in Java, 6th edition**, by M. T. Goodrich, R. Tamassia, and M. H. Goldwasser, Wiley, 2014

Quick-Sort




Quick-Sort

- ◆ Quick-sort is a **randomized** sorting algorithm based on the divide-and-conquer paradigm:
 - **Divide**: pick a **random** element x (called **pivot**) and partition S into
 - ◆ L elements **less** than x
 - ◆ E elements **equal** x
 - ◆ G elements **greater** than x
 - **Recur**: sort L and G
 - **Conquer**: join L , E and G



Partition



- ◆ We partition an input sequence as follows:
 - We remove, in turn, each element y from S and
 - We insert y into L , E or G , depending on the result of the **comparison with the pivot x**
- ◆ Each insertion and removal is at the beginning or at the end of a sequence, and hence takes $O(1)$ time
- ◆ Thus, the partition step of quick-sort takes $O(n)$ time

Algorithm *partition*(S, p)

Input sequence S , position p of pivot

Output subsequences L , E , G of the elements of S less than, equal to, or greater than the pivot, resp.

$L, E, G \leftarrow$ empty sequences

$x \leftarrow S.remove(p)$

while $\neg S.isEmpty()$

$y \leftarrow S.remove(S.first())$

if $y < x$

$L.addLast(y)$

else if $y = x$

$E.addLast(y)$

else $\{ y > x \}$

$G.addLast(y)$


return L, E, G

Java Implementation

```
1  /** Quick-sort contents of a queue. */
2  public static <K> void quickSort(Queue<K> S, Comparator<K> comp) {
3      int n = S.size();
4      if (n < 2) return;                                // queue is trivially sorted
5      // divide
6      K pivot = S.first();                               // using first as arbitrary pivot
7      Queue<K> L = new LinkedList<>();
8      Queue<K> E = new LinkedList<>();
9      Queue<K> G = new LinkedList<>();
10     while (!S.isEmpty()) {                             // divide original into L, E, and G
11         K element = S.dequeue();
12         int c = comp.compare(element, pivot);
13         if (c < 0)                                       // element is less than pivot
14             L.enqueue(element);
15         else if (c == 0)                                 // element is equal to pivot
16             E.enqueue(element);
17         else                                             // element is greater than pivot
18             G.enqueue(element);
19     }
20     // conquer
21     quickSort(L, comp);                                // sort elements less than pivot
22     quickSort(G, comp);                                // sort elements greater than pivot
23     // concatenate results
24     while (!L.isEmpty())
25         S.enqueue(L.dequeue());
26     while (!E.isEmpty())
27         S.enqueue(E.dequeue());
28     while (!G.isEmpty())
29         S.enqueue(G.dequeue());
30 }
```

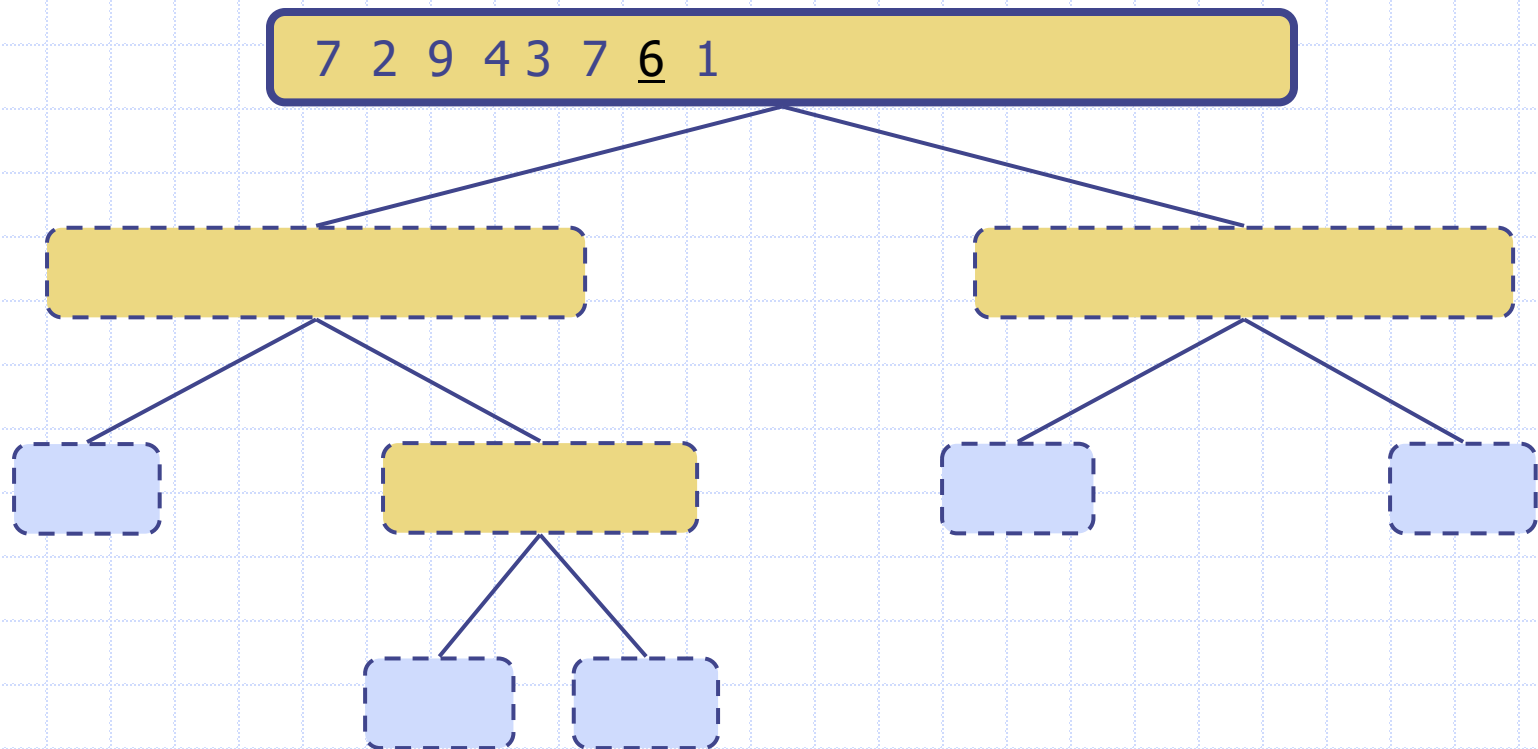
Quick-Sort Tree

- ◆ An execution of quick-sort is depicted by a binary tree
 - Each node represents a recursive call of quick-sort and stores
 - ◆ Unsorted sequence before the execution and its pivot
 - ◆ Sorted sequence at the end of the execution
 - The root is the initial call
 - The leaves are calls on subsequences of size 0 or 1



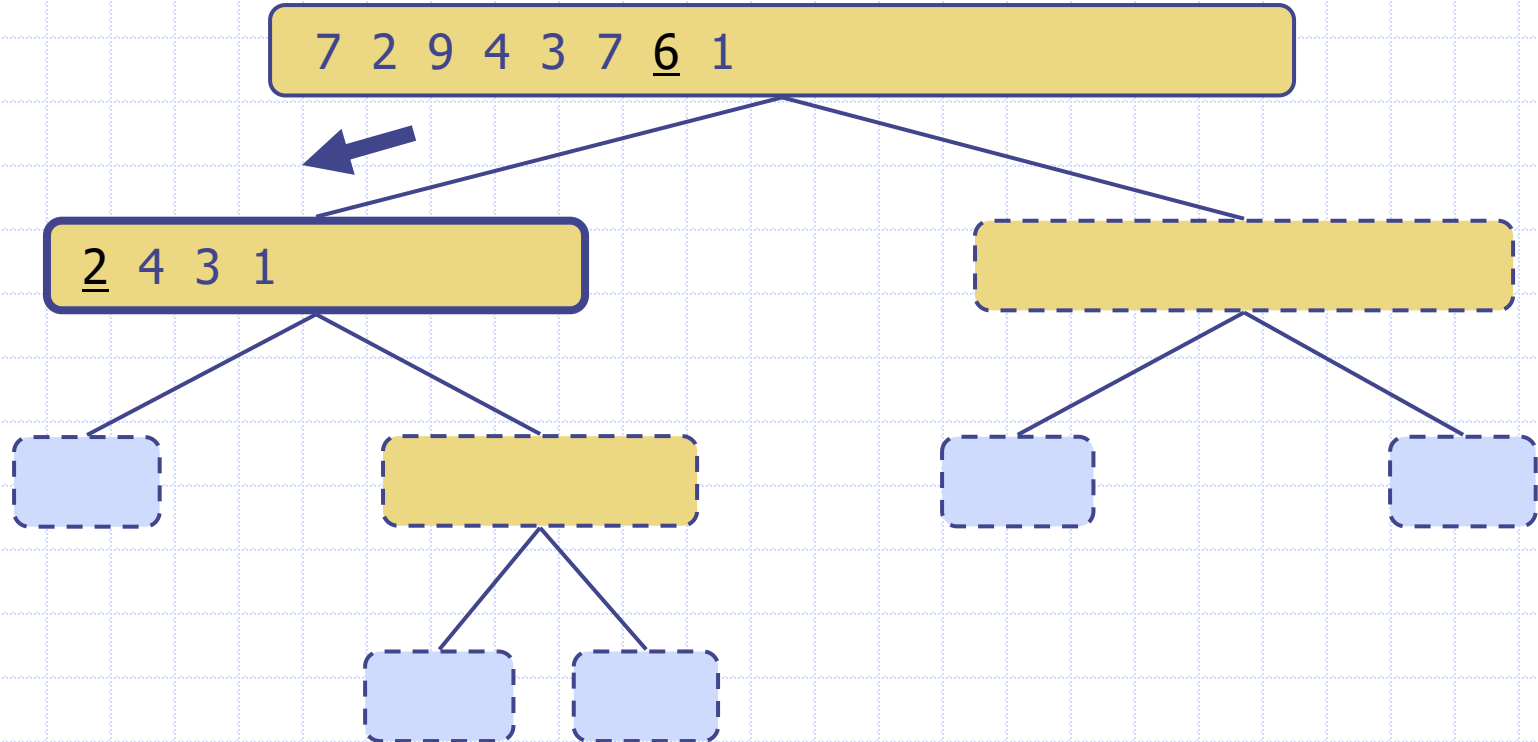
Execution Example

◆ Pivot selection



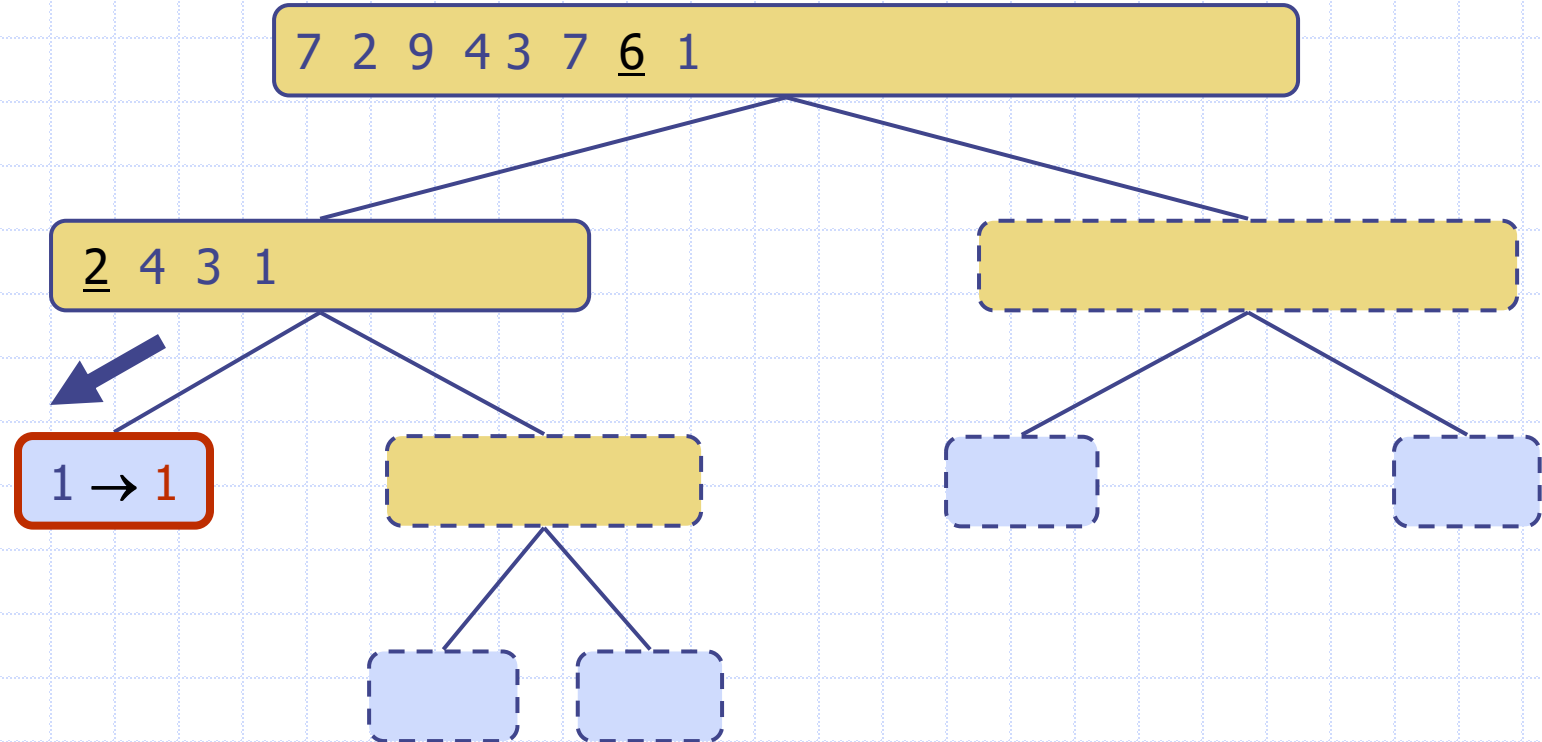
Execution Example (cont.)

- ◆ Partition, recursive call, pivot selection



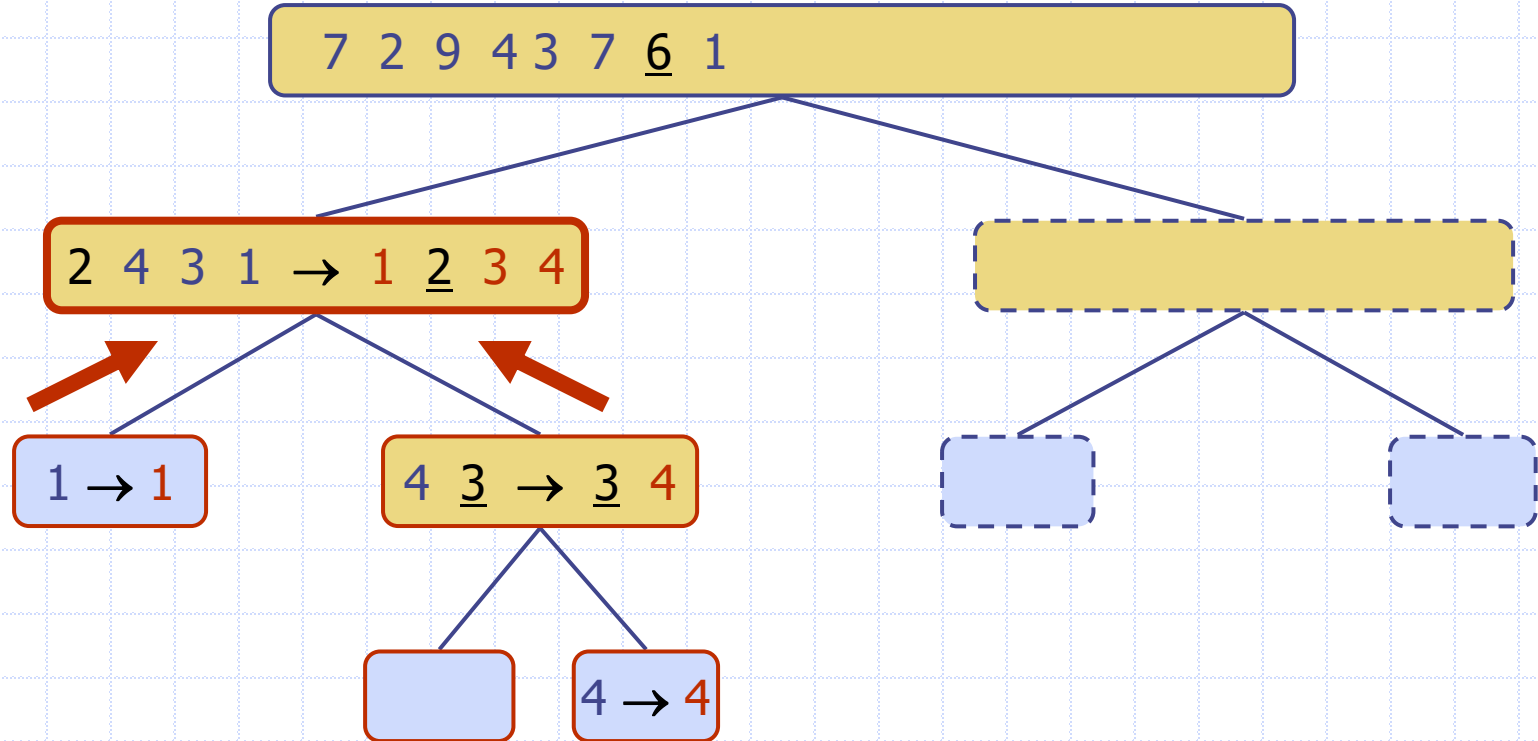
Execution Example (cont.)

- ◆ Partition, recursive call, base case



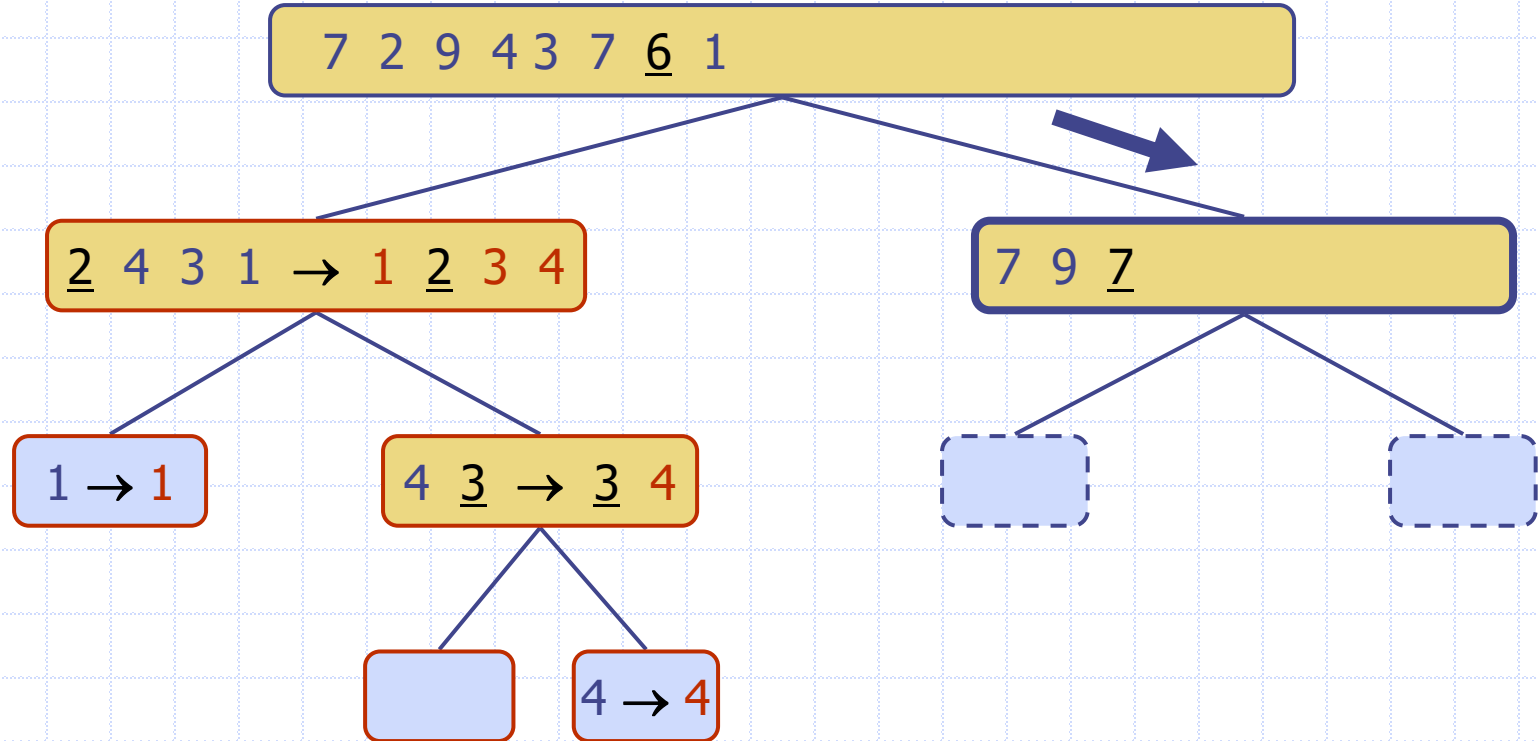
Execution Example (cont.)

- ◆ Recursive call, ..., base case, join



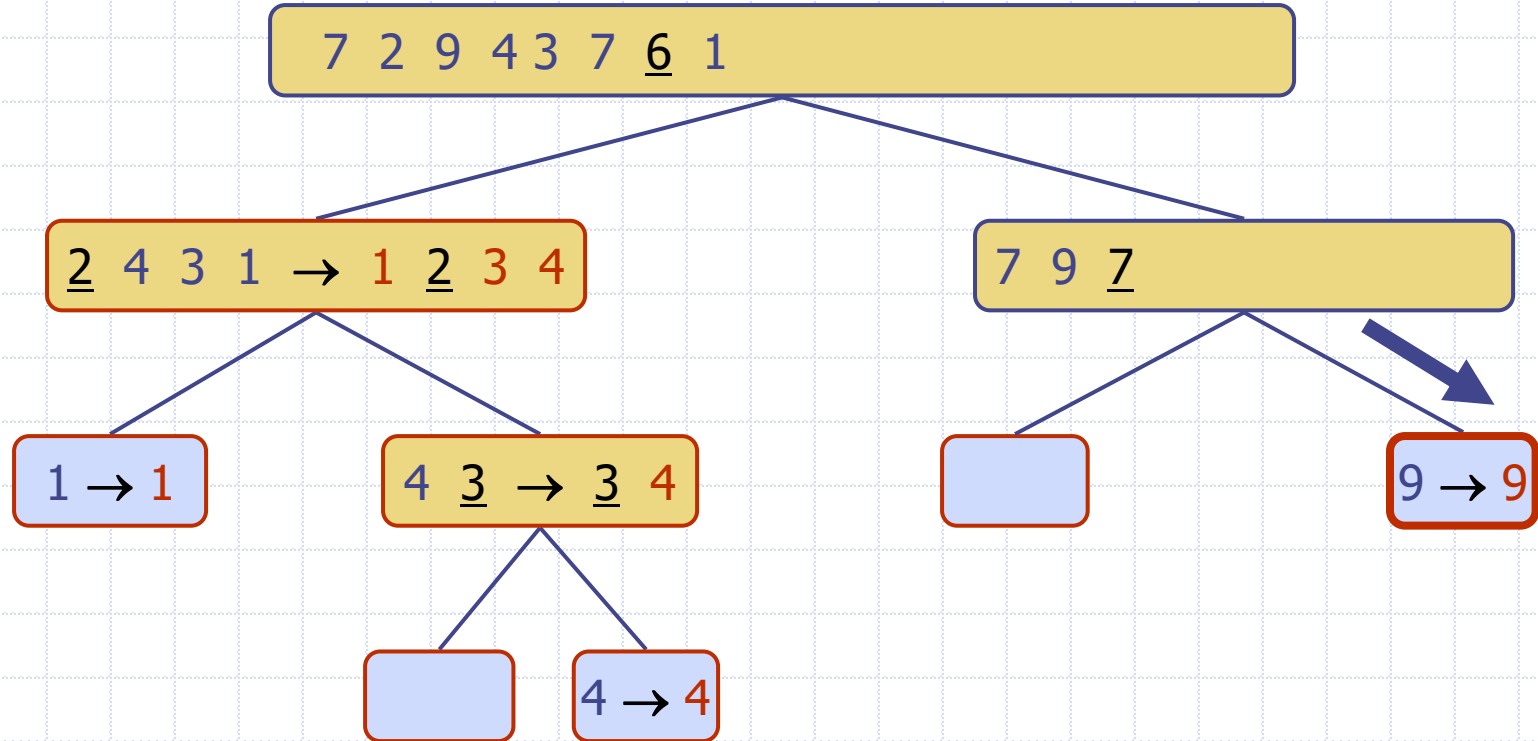
Execution Example (cont.)

- ◆ Recursive call, pivot selection



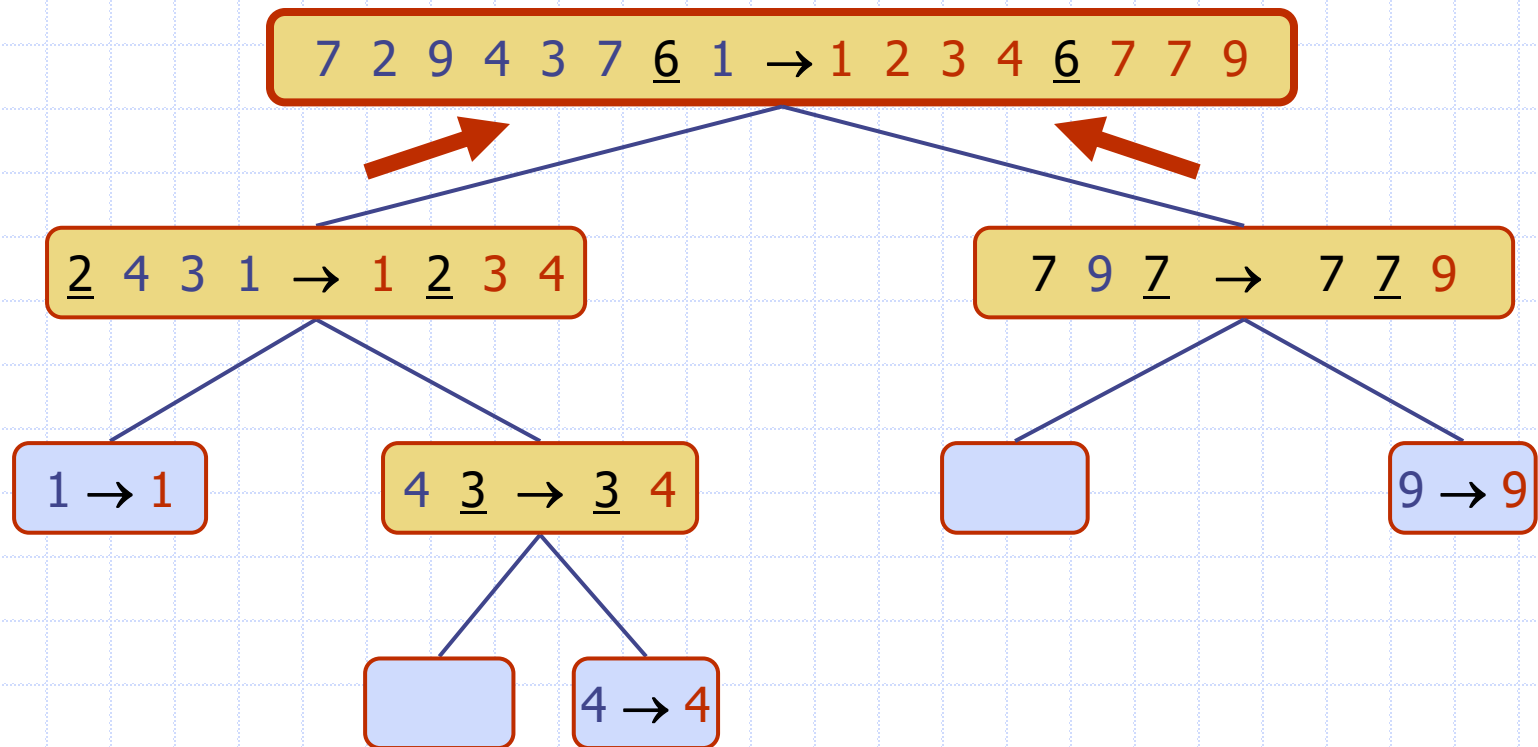
Execution Example (cont.)

- ◆ Partition, ..., recursive call, base case




Execution Example (cont.)

◆ Join, join




Worst-case Running Time

- ◆ The **worst case** for quick-sort occurs when the **pivot is the unique minimum or maximum element**
- ◆ One of L and G has size $n - 1$ and the other has size 0
- ◆ The running time is proportional to the sum
$$n + (n - 1) + \dots + 2 + 1$$
- ◆ Thus, the worst-case running time of quick-sort is $O(n^2)$




Expected Running Time

- ◆ Consider a recursive call of quick-sort on a sequence of size s
 - **Good call:** the sizes of L and G are each less than $3s/4$
 - **Bad call:** one of L and G has size greater than $3s/4$




Good call




Bad call

- ◆ A call is **good** with probability $1/2$
 - $1/2$ of the possible pivots cause good calls:




Expected Running Time, Part 2

- ◆ We can analyze the running time of quick-sort with **the same technique used for merge-sort**.
- ◆ Namely, we can **identify the time spent at each node** of the quick-sort tree T and sum up the running times for all the nodes.
- ◆ The **expected height** of the quick-sort tree is $O(\log n)$
- ◆ The amount of work done at the nodes of the same depth is $O(n)$
- ◆ Thus, the **expected running time of quick-sort is $O(n \log n)$**
- ◆ With a more rigorous analysis, it can be shown that the running time of randomized quick-sort is **$O(n \log n)$ with high probability.**



In-Place Quick-Sort



- ◆ Quick-sort can be implemented to run in-place
- ◆ In the partition step, we use replace operations to rearrange the elements of the input sequence such that
 - the elements less than the pivot have rank less than h
 - the elements equal to the pivot have rank between h and k
 - the elements greater than the pivot have rank greater than k
- ◆ The recursive calls consider
 - elements with rank less than h
 - elements with rank greater than k

Algorithm *inPlaceQuickSort*(S, l, r)

Input sequence S , ranks l and r

Output sequence S with the elements of rank between l and r rearranged in increasing order

if $l \geq r$

return

$i \leftarrow$ a random integer between l and r


$x \leftarrow S.\text{elemAtRank}(i)$

$(h, k) \leftarrow \text{inPlacePartition}(x)$


inPlaceQuickSort($S, l, h - 1$)

inPlaceQuickSort($S, k + 1, r$)


In-Place Partitioning



- ◆ Perform the partition using two indices to split S into L and $E \cup G$ (a similar method can split $E \cup G$ into E and G).



- ◆ Repeat until j and k cross:
 - Scan j to the right until finding an element $\geq x$.
 - Scan k to the left until finding an element $< x$.
 - Swap elements at indices j and k



Java Implementation

```
1  /** Sort the subarray S[a..b] inclusive. */
2  private static <K> void quickSortInPlace(K[ ] S, Comparator<K> comp,
3                                          int a, int b) {
4      if (a >= b) return;          // subarray is trivially sorted
5      int left = a;
6      int right = b-1;
7      K pivot = S[b];
8      K temp;                      // temp object used for swapping
9      while (left <= right) {
10         // scan until reaching value equal or larger than pivot (or right marker)
11         while (left <= right && comp.compare(S[left], pivot) < 0) left++;
12         // scan until reaching value equal or smaller than pivot (or left marker)
13         while (left <= right && comp.compare(S[right], pivot) > 0) right--;
14         if (left <= right) {      // indices did not strictly cross
15             // so swap values and shrink range
16             temp = S[left]; S[left] = S[right]; S[right] = temp;
17             left++; right--;
18         }
19     }
20     // put pivot into its final place (currently marked by left index)
21     temp = S[left]; S[left] = S[b]; S[b] = temp;
22     // make recursive calls
23     quickSortInPlace(S, comp, a, left - 1);
24     quickSortInPlace(S, comp, left + 1, b);
25 }
```

Summary of Sorting Algorithms

Algorithm	Time	Notes
selection-sort	$O(n^2)$	<ul style="list-style-type: none">▪ in-place▪ slow (good for small inputs)
insertion-sort	$O(n^2)$	<ul style="list-style-type: none">▪ in-place▪ slow (good for small inputs)
quick-sort	$O(n \log n)$ expected	<ul style="list-style-type: none">▪ in-place, randomized▪ fastest (good for large inputs)
heap-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ in-place▪ fast (good for large inputs)
merge-sort	$O(n \log n)$	<ul style="list-style-type: none">▪ sequential data access▪ fast (good for huge inputs)